

```
# -*- coding: utf-8 -*-
"""Copy of ENet - Real Time Semantic Segmentation.ipynb
```

Automatically generated by Colaboratory.

Original file is located at
https://colab.research.google.com/drive/1Ar25dVNhaWtRh9hyoo4yOjk_je5_ZNfm

ENet - Real Time Semantic Segmentation

In this notebook, we have reproduced the ENet paper.

 Link to the paper: <https://arxiv.org/pdf/1606.02147.pdf>

 Link to the repository: <https://github.com/iArunava/ENet-Real-Time-Semantic-Segmentation>

Star and Fork!

****ALL THE CODE IN THIS NOTEBOOK ASSUMES THE USAGE OF THE CAMVID DATASET****

```
## Install the dependencies and Import them
"""
```

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import numpy as np
import matplotlib.pyplot as plt
from torch.optim.lr_scheduler import StepLR
import cv2
import os
from tqdm import tqdm
from PIL import Image
```

```
"""## Download the CamVid dataset"""
```

```
!wget https://www.dropbox.com/s/pxcz2wdz04zxocq/CamVid.zip?dl=1 -O CamVid.zip
!unzip CamVid.zip
```

```
"""## Create the ENet model
```

We decided to to split the model to three sub classes:

- 1) Initial block
 - 2) RDDNeck - class for regular, downsampling and dilated bottlenecks
 - 3) ASNeck - class for asymmetric bottlenecks
 - 4) UBNeck - class for upsampling bottlenecks
- ```
"""
```

```
class InitialBlock(nn.Module):
```

```
 # Initial block of the model:
 #
 # Input
 # / \
 # / \
 #maxpool2d conv2d-3x3
 # \ /
 # \ /
 # concatenate
```

```
 def __init__(self, in_channels = 3, out_channels = 13):
 super().__init__()
```

[illegible]

```
 kernel_size = 1)

self.dropout = nn.Dropout2d(p=0.1)

self.conv1 = nn.ConvTranspose2d(in_channels = self.in_channels,
 out_channels = self.reduced_depth,
 kernel_size = 1,
 padding = 0,
 bias = False)

self.prelu1 = activation

This layer used for Upsampling
self.conv2 = nn.ConvTranspose2d(in_channels = self.reduced_depth,
 out_channels = self.reduced_depth,
 kernel_size = 3,
 stride = 2,
 padding = 1,
 output_padding = 1,
 bias = False)

self.prelu2 = activation

self.conv3 = nn.ConvTranspose2d(in_channels = self.reduced_depth,
 out_channels = self.out_channels,
 kernel_size = 1,
 padding = 0,
 bias = False)

self.prelu3 = activation

self.batchnorm = nn.BatchNorm2d(self.reduced_depth)
self.batchnorm2 = nn.BatchNorm2d(self.out_channels)

def forward(self, x, indices):
 x_copy = x

 # Side Branch
 x = self.conv1(x)
 x = self.batchnorm(x)
 x = self.prelu1(x)

 x = self.conv2(x)
 x = self.batchnorm(x)
 x = self.prelu2(x)

 x = self.conv3(x)
 x = self.batchnorm2(x)

 x = self.dropout(x)

 # Main Branch

 x_copy = self.main_conv(x_copy)
 x_copy = self.unpool(x_copy, indices, output_size=x.size())

 # summing the main and side branches
 x = x + x_copy
 x = self.prelu3(x)

 return x

class RDDNeck(nn.Module):
 def __init__(self, dilation, in_channels, out_channels, down_flag, relu=False, projecti
on_ratio=4, p=0.1):
```

```

Regular|Dilated|Downsampling bottlenecks:
#
Bottleneck Input
/ \
/ \
maxpooling2d conv2d-1x1
| | PReLU
| conv2d-3x3
| | PReLU
| conv2d-1x1
| |
Padding2d Regularizer
\ /
Summing + PReLU
#
Params:
dilation (bool) - if True: creating dilation bottleneck
down_flag (bool) - if True: creating downsampling bottleneck
projection_ratio - ratio between input and output channels
relu - if True: relu used as the activation function else: Prelu us used
p - dropout ratio

super().__init__()

Define class variables
self.in_channels = in_channels

self.out_channels = out_channels
self.dilation = dilation
self.down_flag = down_flag

calculating the number of reduced channels
if down_flag:
 self.stride = 2
 self.reduced_depth = int(in_channels // projection_ratio)
else:
 self.stride = 1
 self.reduced_depth = int(out_channels // projection_ratio)

if relu:
 activation = nn.ReLU()
else:
 activation = nn.PReLU()

self.maxpool = nn.MaxPool2d(kernel_size = 2,
 stride = 2,
 padding = 0, return_indices=True)

self.dropout = nn.Dropout2d(p=p)

self.conv1 = nn.Conv2d(in_channels = self.in_channels,
 out_channels = self.reduced_depth,
 kernel_size = 1,
 stride = 1,
 padding = 0,
 bias = False,
 dilation = 1)

self.prelu1 = activation

self.conv2 = nn.Conv2d(in_channels = self.reduced_depth,
 out_channels = self.reduced_depth,
 kernel_size = 3,
 stride = self.stride,
 padding = self.dilation,

```

```

 bias = True,
 dilation = self.dilation)

 self.prelu2 = activation

 self.conv3 = nn.Conv2d(in_channels = self.reduced_depth,
 out_channels = self.out_channels,
 kernel_size = 1,
 stride = 1,
 padding = 0,
 bias = False,
 dilation = 1)

 self.prelu3 = activation

 self.batchnorm = nn.BatchNorm2d(self.reduced_depth)
 self.batchnorm2 = nn.BatchNorm2d(self.out_channels)

def forward(self, x):

 bs = x.size()[0]
 x_copy = x

 # Side Branch
 x = self.conv1(x)
 x = self.batchnorm(x)
 x = self.prelu1(x)

 x = self.conv2(x)
 x = self.batchnorm(x)
 x = self.prelu2(x)

 x = self.conv3(x)
 x = self.batchnorm2(x)

 x = self.dropout(x)

 # Main Branch
 if self.down_flag:
 x_copy, indices = self.maxpool(x_copy)

 if self.in_channels != self.out_channels:
 out_shape = self.out_channels - self.in_channels

 #padding and concatenating in order to match the channels axis of the side and
main branches
 extras = torch.zeros((bs, out_shape, x.shape[2], x.shape[3]))
 if torch.cuda.is_available():
 extras = extras.cuda()
 x_copy = torch.cat((x_copy, extras), dim = 1)

 # Summing main and side branches
 x = x + x_copy
 x = self.prelu3(x)

 if self.down_flag:
 return x, indices
 else:
 return x

class ASNeck(nn.Module):
 def __init__(self, in_channels, out_channels, projection_ratio=4):

 # Asymetric bottleneck:
 #
 # Bottleneck Input
 # / \

```

```
/ \
| | conv2d-1x1
| | PReLU
| | conv2d-1x5
| | |
| | conv2d-5x1
| | |
| | PReLU
| | conv2d-1x1
| | |
Padding2d Regularizer
\ /
Summing + PReLU
#
Params:
projection_ratio - ratio between input and output channels

super().__init__()

Define class variables
self.in_channels = in_channels
self.reduced_depth = int(in_channels / projection_ratio)
self.out_channels = out_channels

self.dropout = nn.Dropout2d(p=0.1)

self.conv1 = nn.Conv2d(in_channels = self.in_channels,
 out_channels = self.reduced_depth,
 kernel_size = 1,
 stride = 1,
 padding = 0,
 bias = False)

self.prelu1 = nn.PReLU()

self.conv21 = nn.Conv2d(in_channels = self.reduced_depth,
 out_channels = self.reduced_depth,
 kernel_size = (1, 5),
 stride = 1,
 padding = (0, 2),
 bias = False)

self.conv22 = nn.Conv2d(in_channels = self.reduced_depth,
 out_channels = self.reduced_depth,
 kernel_size = (5, 1),
 stride = 1,
 padding = (2, 0),
 bias = False)

self.prelu2 = nn.PReLU()

self.conv3 = nn.Conv2d(in_channels = self.reduced_depth,
 out_channels = self.out_channels,
 kernel_size = 1,
 stride = 1,
 padding = 0,
 bias = False)

self.prelu3 = nn.PReLU()

self.batchnorm = nn.BatchNorm2d(self.reduced_depth)
self.batchnorm2 = nn.BatchNorm2d(self.out_channels)

def forward(self, x):
 bs = x.size()[0]
 x_copy = x

 # Side Branch
```

```
x = self.conv1(x)
x = self.batchnorm(x)
x = self.prelu1(x)

x = self.conv21(x)
x = self.conv22(x)
x = self.batchnorm(x)
x = self.prelu2(x)

x = self.conv3(x)

x = self.dropout(x)
x = self.batchnorm2(x)

Main Branch

if self.in_channels != self.out_channels:
 out_shape = self.out_channels - self.in_channels

 #padding and concatenating in order to match the channels axis of the side and
main branches
 extras = torch.zeros((bs, out_shape, x.shape[2], x.shape[3]))
 if torch.cuda.is_available():
 extras = extras.cuda()
 x_copy = torch.cat((x_copy, extras), dim = 1)

Summing main and side branches
x = x + x_copy
x = self.prelu3(x)

return x

class ENet(nn.Module):

 # Creating Enet model!

 def __init__(self, C):
 super().__init__()

 # Define class variables
 # C - number of classes
 self.C = C

 # The initial block
 self.init = InitialBlock()

 # The first bottleneck
 self.b10 = RDDNeck(dilation=1,
 in_channels=16,
 out_channels=64,
 down_flag=True,
 p=0.01)

 self.b11 = RDDNeck(dilation=1,
 in_channels=64,
 out_channels=64,
 down_flag=False,
 p=0.01)

 self.b12 = RDDNeck(dilation=1,
 in_channels=64,
 out_channels=64,
 down_flag=False,
 p=0.01)

 self.b13 = RDDNeck(dilation=1,
 in_channels=64,
```

```
 out_channels=64,
 down_flag=False,
 p=0.01)

self.b14 = RDDNeck(dilation=1,
 in_channels=64,
 out_channels=64,
 down_flag=False,
 p=0.01)

The second bottleneck
self.b20 = RDDNeck(dilation=1,
 in_channels=64,
 out_channels=128,
 down_flag=True)

self.b21 = RDDNeck(dilation=1,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b22 = RDDNeck(dilation=2,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b23 = ASNeck(in_channels=128,
 out_channels=128)

self.b24 = RDDNeck(dilation=4,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b25 = RDDNeck(dilation=1,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b26 = RDDNeck(dilation=8,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b27 = ASNeck(in_channels=128,
 out_channels=128)

self.b28 = RDDNeck(dilation=16,
 in_channels=128,
 out_channels=128,
 down_flag=False)

The third bottleneck
self.b31 = RDDNeck(dilation=1,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b32 = RDDNeck(dilation=2,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b33 = ASNeck(in_channels=128,
 out_channels=128)
```



```
self.b34 = RDDNeck(dilation=4,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b35 = RDDNeck(dilation=1,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b36 = RDDNeck(dilation=8,
 in_channels=128,
 out_channels=128,
 down_flag=False)

self.b37 = ASNeck(in_channels=128,
 out_channels=128)

self.b38 = RDDNeck(dilation=16,
 in_channels=128,
 out_channels=128,
 down_flag=False)

The fourth bottleneck
self.b40 = UBNeck(in_channels=128,
 out_channels=64,
 relu=True)

self.b41 = RDDNeck(dilation=1,
 in_channels=64,
 out_channels=64,
 down_flag=False,
 relu=True)

self.b42 = RDDNeck(dilation=1,
 in_channels=64,
 out_channels=64,
 down_flag=False,
 relu=True)

The fifth bottleneck
self.b50 = UBNeck(in_channels=64,
 out_channels=16,
 relu=True)

self.b51 = RDDNeck(dilation=1,
 in_channels=16,
 out_channels=16,
 down_flag=False,
 relu=True)

Final ConvTranspose Layer
self.fullconv = nn.ConvTranspose2d(in_channels=16,
 out_channels=self.C,
 kernel_size=3,
 stride=2,
 padding=1,
 output_padding=1,
 bias=False)
```

```
def forward(self, x):
```

```
The initial block
x = self.init(x)
```

```
The first bottleneck
x, i1 = self.b10(x)
x = self.b11(x)
x = self.b12(x)
x = self.b13(x)
x = self.b14(x)

The second bottleneck
x, i2 = self.b20(x)
x = self.b21(x)
x = self.b22(x)
x = self.b23(x)
x = self.b24(x)
x = self.b25(x)
x = self.b26(x)
x = self.b27(x)
x = self.b28(x)

The third bottleneck
x = self.b31(x)
x = self.b32(x)
x = self.b33(x)
x = self.b34(x)
x = self.b35(x)
x = self.b36(x)
x = self.b37(x)
x = self.b38(x)

The fourth bottleneck
x = self.b40(x, i2)
x = self.b41(x)
x = self.b42(x)

The fifth bottleneck
x = self.b50(x, i1)
x = self.b51(x)

Final ConvTranspose Layer
x = self.fullconv(x)

return x

"""## Instantiate the ENet model"""

enet = ENet(12)

"""Move the model to cuda if available"""

Checking if there is any gpu available and pass the model to gpu or cpu
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
enet = enet.to(device)

"""## Define the loader that will load the input and output images"""

def loader(training_path, segmented_path, batch_size, h=320, w=1000):
 filenames_t = os.listdir(training_path)
 total_files_t = len(filenames_t)

 filenames_s = os.listdir(segmented_path)
 total_files_s = len(filenames_s)

 assert(total_files_t == total_files_s)

 if str(batch_size).lower() == 'all':
 batch_size = total_files_s

 idx = 0
```

```
while(1):
 # Choosing random indexes of images and labels
 batch_idx = np.random.randint(0, total_files_s, batch_size)

 inputs = []
 labels = []

 for jj in batch_idx:
 # Reading normalized photo
 img = plt.imread(training_path + filenames_t[jj])
 # Resizing using nearest neighbor method
 img = cv2.resize(img, (h, w), cv2.INTER_NEAREST)
 inputs.append(img)

 # Reading semantic image
 img = Image.open(segmented_path + filenames_s[jj])
 img = np.array(img)
 # Resizing using nearest neighbor method
 img = cv2.resize(img, (h, w), cv2.INTER_NEAREST)
 labels.append(img)

 inputs = np.stack(inputs, axis=2)
 # Changing image format to C x H x W
 inputs = torch.tensor(inputs).transpose(0, 2).transpose(1, 3)

 labels = torch.tensor(labels)

 yield inputs, labels

"""## Define the class weights"""

def get_class_weights(num_classes, c=1.02):
 pipe = loader('/content/train/', '/content/trainannot/', batch_size='all')
 _, labels = next(pipe)
 all_labels = labels.flatten()
 each_class = np.bincount(all_labels, minlength=num_classes)
 propensity_score = each_class / len(all_labels)
 class_weights = 1 / (np.log(c + propensity_score))
 return class_weights

class_weights = get_class_weights(12)

"""## Define the Hyperparameters"""

lr = 5e-4
batch_size = 10

criterion = nn.CrossEntropyLoss(weight=torch.FloatTensor(class_weights).to(device))
optimizer = torch.optim.Adam(enet.parameters(),
 lr=lr,
 weight_decay=2e-4)

print_every = 5
eval_every = 5

"""## Training loop"""

train_losses = []
eval_losses = []

bc_train = 367 // batch_size # mini_batch train
bc_eval = 101 // batch_size # mini_batch validation

Define pipeline objects
pipe = loader('/content/train/', '/content/trainannot/', batch_size)
eval_pipe = loader('/content/val/', '/content/valannot/', batch_size)
```

```
epochs = 100

Train loop
for e in range(1, epochs+1):

 train_loss = 0
 print ('-'*15, 'Epoch %d' % e, '-'*15)

 enet.train()

 for _ in tqdm(range(bc_train)):
 X_batch, mask_batch = next(pipe)

 # assign data to cpu/gpu
 X_batch, mask_batch = X_batch.to(device), mask_batch.to(device)

 optimizer.zero_grad()

 out = enet(X_batch.float())

 # loss calculation
 loss = criterion(out, mask_batch.long())
 # update weights
 loss.backward()
 optimizer.step()

 train_loss += loss.item()

 print ()
 train_losses.append(train_loss)

 if (e+1) % print_every == 0:
 print ('Epoch {}/{}...'.format(e, epochs),
 'Loss {:.6f}'.format(train_loss))

 if e % eval_every == 0:
 with torch.no_grad():
 enet.eval()

 eval_loss = 0

 # Validation loop
 for _ in tqdm(range(bc_eval)):
 inputs, labels = next(eval_pipe)

 inputs, labels = inputs.to(device), labels.to(device)

 out = enet(inputs)

 out = out.data.max(1)[1]

 eval_loss += (labels.long() - out.long()).sum()

 print ()
 print ('Loss {:.6f}'.format(eval_loss))

 eval_losses.append(eval_loss)

 if e % print_every == 0:
 checkpoint = {
 'epochs' : e,
 'state_dict' : enet.state_dict()
```

```

 }
 torch.save(checkpoint, '/content/ckpt-enet-{}-{}.pth'.format(e, train_loss))
 print ('Model saved!')

print ('Epoch {}/{}...'.format(e, epochs),
 'Total Mean Loss: {:.6f}'.format(sum(train_losses) / epochs))

"""## Infer using the trained model

Get the PreTrained ENet model
"""

!wget https://github.com/iArunava/ENet-Real-Time-Semantic-Segmentation/raw/master/datasets/
CamVid/ckpt-enet.pth

"""Load the ENet model"""

Load a pretrained model if needed
enet = ENet(12)
state_dict = torch.load('/content/ckpt-enet.pth')['state_dict']
enet.load_state_dict(state_dict)

"""## Use the code to infer on new images"""

fname = 'Seq05VD_f05100.png'
tmg_ = plt.imread('/content/test/' + fname)
tmg_ = cv2.resize(tmg_, (512, 512), cv2.INTER_NEAREST)
tmg = torch.tensor(tmg_).unsqueeze(0).float()
tmg = tmg.transpose(2, 3).transpose(1, 2).to(device)

enet.to(device)
with torch.no_grad():
 out1 = enet(tmg.float()).squeeze(0)

"""## Load the label image"""

smg_ = Image.open('/content/testannot/' + fname)
smg_ = cv2.resize(np.array(smg_), (512, 512), cv2.INTER_NEAREST)

"""## Move the output to cpu and convert it to numpy and see how each class looks"""

out2 = out1.cpu().detach().numpy()

mno = 8 # Should be between 0 - n-1 | where n is the number of classes

figure = plt.figure(figsize=(20, 10))
plt.subplot(1, 3, 1)
plt.title('Input Image')
plt.axis('off')
plt.imshow(tmg_)
plt.subplot(1, 3, 2)
plt.title('Output Image')
plt.axis('off')
plt.imshow(out2[mno, :, :])
plt.show()

"""Get the class labels from the output"""

b_ = out1.data.max(0)[1].cpu().numpy()

"""Define the function that maps a 2D image with all the class labels to a segmented image
with the specified colored maps"""

def decode_segmap(image):
 Sky = [128, 128, 128]
 Building = [128, 0, 0]
 Pole = [192, 192, 128]
 Road_marking = [255, 69, 0]

```

```
Road = [128, 64, 128]
Pavement = [60, 40, 222]
Tree = [128, 128, 0]
SignSymbol = [192, 128, 128]
Fence = [64, 64, 128]
Car = [64, 0, 128]
Pedestrian = [64, 64, 0]
Bicyclist = [0, 128, 192]

label_colours = np.array([Sky, Building, Pole, Road_marking, Road,
 Pavement, Tree, SignSymbol, Fence, Car,
 Pedestrian, Bicyclist]).astype(np.uint8)
r = np.zeros_like(image).astype(np.uint8)
g = np.zeros_like(image).astype(np.uint8)
b = np.zeros_like(image).astype(np.uint8)
for l in range(0, 12):
 r[image == l] = label_colours[l, 0]
 g[image == l] = label_colours[l, 1]
 b[image == l] = label_colours[l, 2]

rgb = np.zeros((image.shape[0], image.shape[1], 3)).astype(np.uint8)
rgb[:, :, 0] = b
rgb[:, :, 1] = g
rgb[:, :, 2] = r
return rgb

"""Decode the images"""

true_seg = decode_segmap(sm_g_)
pred_seg = decode_segmap(b_)

figure = plt.figure(figsize=(20, 10))
plt.subplot(1, 3, 1)
plt.title('Input Image')
plt.axis('off')
plt.imshow(tmg_)
plt.subplot(1, 3, 2)
plt.title('Predicted Segmentation')
plt.axis('off')
plt.imshow(pred_seg)
plt.subplot(1, 3, 3)
plt.title('Ground Truth')
plt.axis('off')
plt.imshow(true_seg)
plt.show()

"""## Save the model checkpoint"""

Save the parameters
checkpoint = {
 'epochs' : e,
 'state_dict' : enet.state_dict()
}
torch.save(checkpoint, 'ckpt-enet-1.pth')

"""## Save the entire model"""

Save the model
torch.save(enet, '/content/model.pt')
```